

CIL Cheat Sheet 2015

Jan Wilken Dörrie

LINEAR ALGEBRA PRIMER

Equivalent Conditions

For $\mathbf{A} \in \mathbb{R}^{M \times M}$ the following conditions are equivalent:
 \mathbf{A} has an inverse \mathbf{A}^{-1} ; $\text{rank}(\mathbf{A}) = M$; $\text{range}(\mathbf{A}) = \mathbb{R}^M$;
 $\text{null}(\mathbf{A}) = \{\mathbf{0}\}$; 0 is not an eigenvalue of \mathbf{A} ; 0 is not a singular value of \mathbf{A} .

NORMS

Vector norms

A *norm* is a function $\|\cdot\| : V \rightarrow \mathbb{R}$ quantifying the size of a vector. It must satisfy

- 1) Positive scalability: $\|a \cdot \mathbf{x}\| = |a| \cdot \|\mathbf{x}\|$ for $a \in \mathbb{R}$
 - 2) Triangle inequality: $\|\mathbf{x} + \mathbf{y}\| \leq \|\mathbf{x}\| + \|\mathbf{y}\|$, $\mathbf{x}, \mathbf{y} \in V$.
 - 3) Separability: $\|\mathbf{x}\| = 0$ implies $\mathbf{x} = \mathbf{0}$.
- Most common are *p-norms*: $\|\mathbf{x}\|_p := (\sum_{i=1}^n |x_i|^p)^{1/p}$
 - Special case is *Euclidean norm*: $\|\mathbf{x}\|_2 := \sqrt{\sum_{i=1}^n x_i^2}$
 - The “0-norm” is $\|\mathbf{x}\|_0 := |\{x_i \mid x_i \neq 0\}|$

Matrix norms

We can also define norms on matrices, satisfying the properties described above. $\mathbf{A} \in \mathbb{R}^{M \times N}$:

- *Frobenius*: $\|\mathbf{A}\|_F := \sqrt{\sum_{ij} a_{ij}^2} = \sqrt{\sum_{i=1}^{\min(M,N)} \sigma_i^2}$
- *p-norms for matrices*: $\|\mathbf{A}\|_p := \sup\{\|\mathbf{A}\mathbf{x}\|_p / \|\mathbf{x}\|_p\}$
- *Euclidean*: $\|\mathbf{A}\|_2 := \sup\{\|\mathbf{A}\mathbf{x}\|_2 / \|\mathbf{x}\|_2\} = \sigma_{\max}$
- *Nuclear norm*: $\|\mathbf{A}\|_* := \sum_{i=1}^{\min(M,N)} \sigma_i$

STATISTICS

Kullback-Leibler Divergence

- Divergence between discrete probability distributions P and Q : $D_{\text{KL}}(P\|Q) = \sum_{\omega \in \Omega} P(\omega) \log \left(\frac{P(\omega)}{Q(\omega)} \right)$.
- Properties of the Kullback-Leibler Divergence
 - $D_{\text{KL}}(P\|Q) \geq 0$.
 - $D_{\text{KL}}(P\|Q) = 0$ if and only if P and Q are identical.
 - $D_{\text{KL}}(P\|Q) \neq D_{\text{KL}}(Q\|P)$!
 - caution: the KL-Divergence is not symmetric, therefore it is not a metric/distance!

DIMENSION REDUCTION

Principal Component Analysis (PCA)

Orthogonal linear projection of high dimensional data onto low dimensional subspace. Objectives:

- 1) Minimize error $\|\mathbf{x} - \tilde{\mathbf{x}}\|_2$ of point \mathbf{x} and its approximation $\tilde{\mathbf{x}}$.
- 2) Preserve information: maximize variance.

Both objectives are shown to be formally equivalent.

Statistics of Projected Data:

- Mean of the data: sample mean $\bar{\mathbf{x}}$
- Covariance of the data:

$$\Sigma = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^\top$$

Solution: Eigenvalue Decomposition: The eigenvalue decomposition of the covariance matrix $\Sigma = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^\top$ contains all relevant information.

- For $K \leq D$ dimensional projection space: Choose K eigenvectors $\{\mathbf{u}_1, \dots, \mathbf{u}_K\}$ with largest associated eigenvalues $\{\lambda_1, \dots, \lambda_K\}$.

Singular Value Decomposition

Theorem (Eckart-Young). Let \mathbf{A} be a matrix of rank R , if we wish to approximate \mathbf{A} using a matrix of a lower rank K then, $\tilde{\mathbf{A}} = \sum_{k=1}^K d_k \mathbf{u}_k \mathbf{v}_k^\top$ is the closest matrix in the Frobenius norm. (Assumes ordering of singular values $d_k \geq d_{k+1}$)

CLUSTERING

K-Means

Motivation:

- Given: set of data points $\mathbf{x}_1, \dots, \mathbf{x}_N \in \mathbb{R}^D$
- Goal: find *meaningful partition* of the data
 - i.e. a labeling of each data point with a unique label

$$\pi : \{1, \dots, N\} \rightarrow \{1, \dots, K\} \text{ or } \pi : \mathbb{R}^D \rightarrow \{1, \dots, K\}$$

- note: numbering of clusters is arbitrary
- k -th cluster recovered by $\pi^{-1}(k) \subseteq \{1, \dots, N\}$ or $\subseteq \mathbb{R}^D$

Vector Quantization:

- Partition of the space \mathbb{R}^D
- Clusters represented by *centroids* $\mathbf{u}_k \in \mathbb{R}^D$
- Mapping induced via nearest centroid rule

$$\pi(\mathbf{x}) = \underset{k=1, \dots, K}{\operatorname{argmin}} \|\mathbf{u}_k - \mathbf{x}\|_2$$

Objective Function for K-Means:

- Useful notation: represent π via indicator matrix \mathbf{Z} :

$$z_{kn} := [\pi(\mathbf{x}_n) = k]$$

- K-means Objective function

$$J(\mathbf{U}, \mathbf{Z}) = \sum_{n=1}^N \sum_{k=1}^K z_{kn} \|\mathbf{x}_n - \mathbf{u}_k\|_2^2 = \|\mathbf{X} - \mathbf{U}\mathbf{Z}\|_F^2,$$

where $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_N] \in \mathbb{R}^{D \times N}$ and $\mathbf{U} = [\mathbf{u}_1, \dots, \mathbf{u}_K] \in \mathbb{R}^{D \times K}$

K-means Algorithm: Optimal Assignment:

- Compute optimal assignment \mathbf{Z} , given centroids \mathbf{U}
 - minimize each column of \mathbf{Z} separately

$$\mathbf{z}_{\bullet n}^* = \underset{z_{1n}, \dots, z_{Kn}}{\operatorname{argmin}} \sum_{k=1}^K z_{kn} \|\mathbf{x}_n - \mathbf{u}_k\|_2^2$$

- optimum is attained by mapping to the closest centroid

$$z_{kn}^*(\mathbf{U}) = \left[k = \underset{l}{\operatorname{argmin}} \|\mathbf{x}_n - \mathbf{u}_l\|_2 \right]$$

K-means Algorithm: Optimal Assignment:

- Compute optimal choice of \mathbf{U} , given assignments \mathbf{Z}
 - continuous variables: compute gradient and set to zero (necessary optimality condition)
 - look at (partial) gradient for every centroid \mathbf{u}_k

$$\nabla_{\mathbf{u}_k} J(\mathbf{U}, \mathbf{Z}) = \sum_{n=1}^N z_{kn} \nabla_{\mathbf{u}_k} \|\mathbf{x}_n - \mathbf{u}_k\|_2^2 = -2 \sum_{n=1}^N z_{kn} (\mathbf{x}_n - \mathbf{u}_k)$$

- setting gradient to zero

$$\nabla_{\mathbf{U}} J(\mathbf{U}, \mathbf{Z}) \stackrel{!}{=} 0 \implies \mathbf{u}_k^*(\mathbf{Z}) = \frac{\sum_{n=1}^N z_{kn} \mathbf{x}_n}{\sum_{n=1}^N z_{kn}}$$

K-means Algorithm: Analysis:

- Computational cost of each iteration is $O(KND)$
- K -means convergence is guaranteed
- K -means optimizes a non-convex objective. Hence we are not guaranteed to find the global optimum.
- Finds a local optimum (\mathbf{U}, \mathbf{Z}) in the following sense
 - for each \mathbf{Z}' with $\frac{1}{2} \|\mathbf{Z} - \mathbf{Z}'\|_0 = 1$ (differs in one assignment)
 - $J(\mathbf{U}^*(\mathbf{Z}'), \mathbf{Z}') \geq J(\mathbf{U}, \mathbf{Z})$
 - may gain by changing assignments of ≥ 2 points
- K -means algorithm can be used to compress data
 - with information loss, if $K < N$
 - store only the centroids and the assignments

Mixture Models

Gaussian Mixture Model (GMM):

$$p_\theta(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k),$$

where $\pi_k \geq 0$, $\sum_{k=1}^K \pi_k = 1$.

Complete Data Distribution:

- Explicitly introduce latent variables in the generative model
- Assignment variable (for a generic data point) $z_k \in \{0, 1\}$, $\sum_{k=1}^K z_k = 1$
- We have that $\Pr(z_k = 1) = \pi_k$ or $p(\mathbf{z}) = \prod_{k=1}^K \pi_k^{z_k}$
- Joint distribution over (\mathbf{x}, \mathbf{z}) (*complete data* distribution) $p(\mathbf{x}, \mathbf{z}) = \prod_{k=1}^K [\pi_k \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)]^{z_k}$

Posterior Assignments: Posterior probabilities for assignments

$$\Pr(z_k = 1 \mid \mathbf{x}) = \frac{\pi_k \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{l=1}^K \pi_l \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}_l, \boldsymbol{\Sigma}_l)}$$

Lower Bounding the Log-Likelihood:

- Expectation Maximization
 - maximize a lower bound on the log-likelihood
 - systematic way of deriving a family of bounds
 - based on complete data distribution

- Specifically:

$$\begin{aligned} \ln p_\theta(\mathbf{x}) &= \ln \sum_{\mathbf{z}} p_\theta(\mathbf{x}, \mathbf{z}) = \ln \sum_{k=1}^K p(\mathbf{x}, \theta_k) \pi_k \\ &= \ln \sum_{k=1}^K q_k \frac{p(\mathbf{x}; \theta_k) \pi_k}{q_k} \\ &\geq \sum_{k=1}^K q_k [\ln p(\mathbf{x}, \theta_k) + \ln \pi_k - \ln q_k] \end{aligned}$$

- follows from Jensen's inequality (concavity of logarithm)
- can be done for the contribution of each data point (additive)

Mixture Model: Expectation Step:

$$q_k = \frac{\pi_k p(\mathbf{x}; \theta_k)}{\sum_{l=1}^K \pi_l p(\mathbf{x}, \theta_l)} = \Pr(z_k = 1 \mid \mathbf{x})$$

Mixture Model: Maximization Step:

$$\begin{aligned} \pi_k^* &= \frac{1}{N} \sum_{n=1}^N q_{kn} \\ \boldsymbol{\mu}_k^* &= \frac{\sum_{n=1}^N q_{kn} \mathbf{x}_n}{\sum_{n=1}^N q_{kn}} \\ \boldsymbol{\Sigma}_k^* &= \frac{\sum_{n=1}^N q_{kn} (\mathbf{x}_n - \boldsymbol{\mu}_k^*)(\mathbf{x}_n - \boldsymbol{\mu}_k^*)^\top}{\sum_{n=1}^N q_{kn}} \end{aligned}$$

AIC and BIC:

- Trade-off: achieve balance between data fit — measured by likelihood $p(\mathbf{X} \mid \theta)$ — and complexity. Complexity can be measured by the number of free parameters $\kappa(\cdot)$.
- Different Heuristics for choosing K
 - Akaike Information Criterion (AIC)
 - Bayesian Information Criterion (BIC)

$$\text{AIC}(\theta \mid \mathbf{X}) = -\ln p_\theta(\mathbf{X}) + \kappa(\theta)$$

$$\text{BIC}(\theta \mid \mathbf{X}) = -\ln p_\theta(\mathbf{X}) + \frac{1}{2} \kappa(\theta) \ln N$$

- Generally speaking, the BIC criterion penalizes complexity more than the AIC criterion.

Non-Negative Matrix Factorization

Non-Negative Matrix Factorization:

- Document-term matrix $\mathbf{X} \in \mathbb{R}_{\geq 0}^{D \times N}$ storing the word counts for each document:

$$\mathbf{X} = \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$$

x_{dn} : Frequency of the d -th word in the n -th document.

- Non-negative matrix factorization (NMF) of \mathbf{X} :

$$\mathbf{X} \approx \mathbf{U} \mathbf{Z}$$

- with $\mathbf{U} \in \mathbb{R}_{\geq 0}^{D \times K}$ and $\mathbf{Z} \in \mathbb{R}_{\geq 0}^{K \times N}$
 - * N : number of documents
 - * D : vocabulary size
 - * K : number of dimensions (design choice)
 - * data reduction: $(D + N)K \ll DN$

pLSI — Generative Model:

- For a given *document* sample $\text{len}(\text{document})$ words by a two-stage procedure:
 - sample a topic according to $P(\text{topic} \mid \text{document})$
 - sample a word according to $P(\text{word} \mid \text{topic})$
- Key assumption: *conditional independence* of word and document given topic
- Conditional distribution of a word, given a document:

$$P(\text{word} \mid \text{document}) = \sum_{k=1}^K P(\text{word} \mid \text{topic}_k) P(\text{topic}_k \mid \text{document}) \implies q_{tk}^* \propto u_{dk} z_{kn_t}, \text{ i.e. } q_{tk}^* = \frac{u_{dk} z_{kn_t}}{\sum_{l=1}^K u_{dl} z_{ln_t}}$$

- Side note: how to sample a “new” document? Can use fully generative model of LDA.

pLSI — Matrix Factorization View:

- *Normalize* the elements of \mathbf{X} so that they correspond to relative frequencies:

$$T := \sum_{d=1}^D \sum_{n=1}^N x_{dn}, \quad x_{dn} \leftarrow \frac{x_{dn}}{T}$$

- *Matrix Factorization*
 - pLSI can be understood as a matrix factorization of the form $\mathbf{X} \approx \mathbf{U}\mathbf{Z}$, with $\mathbf{U} \in \mathbb{R}_{\geq 0}^{D \times K}$, and $\mathbf{Z} \in \mathbb{R}_{\geq 0}^{K \times N}$
 - where additionally we have the constraints:
 - * $\sum_{d=1}^D u_{dk} = 1 (\forall k)$, identify $u_{dk} \equiv P(\text{word}_d \mid \text{topic}_k)$
 - * $\sum_{k,n} z_{kn} = 1$, identify $z_{kn} \equiv P(\text{topic}_k \mid \text{document}_n) P(\text{document}_n)$

pLSI — Parameter Estimation:

- Goal: maximize the likelihood of the data under the model
- Data: the relative frequencies \mathbf{X}
- Probabilistic model: $P(\text{word}_d, \text{document}_n) = \sum_{k=1}^K P(\text{word}_d \mid \text{topic}_k) P(\text{topic}_k \mid \text{document}_n) = (\mathbf{U}\mathbf{Z})_{dn}$
- *Log likelihood*: $\log \mathcal{L}(\mathbf{U}, \mathbf{Z}; \mathbf{X}) = \log P(\mathbf{X}; \mathbf{U}, \mathbf{Z}) = \sum_{d=1}^D \sum_{n=1}^N x_{dn} \log \sum_{k=1}^K u_{dk} z_{kn}$

EM for pLSI — Variational Likelihood:

- Follow similar recipe as for Gaussian Mixture Model
- Reindex the observations in a per token manner with $t = 1, \dots, T$
 - pairs of word/documents indexes (d_t, n_t)
 - note that $\sum_{t=1}^T f(d_t, n_t) = \sum_{d=1}^D \sum_{n=1}^N x_{dn} f(d, n)$ for arbitrary functions f
- *Variational Likelihood*

$$\begin{aligned} \log P(\mathbf{X}; \mathbf{U}, \mathbf{Z}) &= \sum_{t=1}^T \log (\mathbf{U}\mathbf{Z})_{d_t n_t} = \sum_{t=1}^T \log \left[\sum_{k=1}^K u_{d_t k} z_{k n_t} \right] \\ &\geq \sum_{t=1}^T \max_{q \in \mathcal{S}_K} \sum_{k=1}^K q_k [\log u_{d_t k} + \log z_{k n_t} - \log q_k] \end{aligned}$$

- $\mathcal{S}_K := \left\{ x \in \mathbb{R}^K \mid x \geq 0, \sum_{k=1}^K x_k = 1 \right\}$ (probability simplex)

EM for pLSI — Derivation of E-step:

- Compute the argmin in the variational bound

$$q_t^* = \argmax_{q \in \mathcal{S}_K} \sum_{k=1}^K q_k [\log u_{d_t k} + \log z_{k n_t} + \log q_k]$$

- Form Lagrangian and differentiate

$$\frac{\partial}{\partial q_k} \{ q_k [\log u_{d_t k} + \log z_{k n_t} - \log q_k - \lambda_t^*] \} \stackrel{!}{=} 0$$

- $q_{tk}^* =$ posterior probability that t -th token (i.e. word with index d_t in document with index n_t) has been generated from topic k

EM for pLSI — Derivation of M-step:

- Differentiate lower bound with plugged in optimal choices for q_t^* ($t = 1, \dots, T$)
- M-step solution for \mathbf{U} and \mathbf{Z}

$$u_{dk}^* = \frac{\sum_{t: d_t=d} q_{tk}^*}{\sum_{t=1}^T q_{tk}^*} \quad z_{kn}^* = \frac{\sum_{t: n_t=n} q_{tk}^*}{T}$$

SPARSE CODING

Optimization

Coordinate Descent: Idea: Update one coordinate at a time, while keeping others fixed.

- Algorithm:

- initialize $\mathbf{x}^{(0)} \in \mathbb{R}^D$
- for $t = 0, \dots, \text{maxIter}$

$$* \quad d \leftarrow \mathcal{U}\{1, D\}$$

$$* \quad u^* \leftarrow \argmin_{u \in \mathbb{R}} f \left(x_1^{(t)}, \dots, x_{d-1}^{(t)}, u, x_{d+1}^{(t)}, \dots, x_D^{(t)} \right)$$

$$* \quad x_d^{(t+1)} \leftarrow u^*, \quad x_{d'}^{(t+1)} \leftarrow x_{d'}^{(t)} \text{ for } d' \neq d$$

Gradient Descent Method:

- Algorithm:

- initialize $\mathbf{x}^{(0)} \in \mathbb{R}^D$

- for $t = 0, \dots, \text{maxIter}$

$$* \quad \mathbf{x}^{(t+1)} \leftarrow \mathbf{x}^{(t)} - \gamma \nabla f(\mathbf{x}^{(t)})$$

- simple to implement

- good scalability and robustness

- *stepsize* γ usually decreasing with $\gamma \approx \frac{1}{t}$

Stochastic Gradient Descent:

- Optimization Problem Structure: minimize $f(\mathbf{x}) = \frac{1}{N} \sum_{n=1}^N f_n(\mathbf{x})$ with $\mathbf{x} \in \mathbb{R}^D$

- Algorithm:

- initialize $\mathbf{x}^{(0)} \in \mathbb{R}^D$

- for $t = 0, \dots, \text{maxIter}$

$$* \quad n \leftarrow \mathcal{U}\{1, N\}$$

$$* \quad \mathbf{x}^{(t+1)} \leftarrow \mathbf{x}^{(t)} - \gamma \nabla f_n(\mathbf{x}^{(t)})$$

Duality for Constrained Optimization:

- Constrained Problem Formulation (Standard Form): minimize $f(\mathbf{x})$ subject to $g_i(\mathbf{x}) \leq 0, i = 1, \dots, m, h_i(\mathbf{x}) = 0, i = 1, \dots, p$
- Unconstrained Problem: minimize $f(\mathbf{x}) + \sum_{i=1}^m I_-(g_i(\mathbf{x})) + \sum_{i=1}^p I_0(h_i(\mathbf{x}))$. I_- and I_0 are “brickwall” indicator functions.

Dual Problem:

- Lagrangian: $L(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\nu}) := f(\mathbf{x}) + \sum_{i=1}^m \lambda_i g_i(\mathbf{x}) + \sum_{i=1}^p \nu_i h_i(\mathbf{x})$
- Lagrange dual function: $d(\boldsymbol{\lambda}, \boldsymbol{\nu}) := \inf_{\mathbf{x}} L(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\nu})$
- Lagrange *dual problem*: maximize $d(\boldsymbol{\lambda}, \boldsymbol{\nu})$ subject to $\boldsymbol{\lambda} \geq \mathbf{0}$.
 - It is always a lower bound on the primal value $f(\mathbf{x})$ of any feasible \mathbf{x} and thus a lower bound on the unknown solution value $f(\mathbf{x}^*)$ of the primal problem.
 - Strong Duality: If the primal optimization problem is convex and under some additional conditions, the solution value of the dual problem is *equal* to the solution value $f(\mathbf{x}^*)$ of the primal problem.

Convexity:

- Convex Set: A set Q is convex if for any $\mathbf{x}, \mathbf{y} \in Q$ and any $\theta \in [0, 1]$, we have $\theta\mathbf{x} + (1 - \theta)\mathbf{y} \in Q$.
- Convex Function: A function $f: \mathbb{R}^D \rightarrow \mathbb{R}$ is convex if $\text{dom } f$ is a convex set and if for all $\mathbf{x}, \mathbf{y} \in \text{dom } f$, and $\theta \in [0, 1]$ we have $f(\theta\mathbf{x} + (1 - \theta)\mathbf{y}) \leq \theta f(\mathbf{x}) + (1 - \theta)f(\mathbf{y})$.
- Convex Optimization: Convex Optimization Problems are of the form $\min f(\mathbf{x})$ s.t. $\mathbf{x} \in Q$ where both f is a convex function and Q is a convex set (note: \mathbb{R}^D is convex). In Convex Optimization Problems every local minimum is a *global minimum*.

ROBUST PCA